



SYSTEM OVERVIEW · DUE DILIGENCE

Agent Workforce Platform

An operating system for a multi-agent AI workforce.

The Platform is not a chatbot or a prompt library. It is a running operating system that lets an organization command a team of specialized AI agents the way a leader commands a company — with a **chain of command, durable memory, real tool access, human sign-off, and full operational visibility**.

Its core idea: **bring or build your own agents, and get the hard part out of the box**. Standing up one useful agent is easy; making many agents work together reliably, remember what matters, and stay under human control is the hard problem. The Platform solves that hard problem — providing the orchestration, memory, context management, governance, and observability that turn a collection of prompts into a dependable digital workforce.

The value is not any single agent. It is the *system* that makes many agents work together reliably, safely, and without constant babysitting.

3,238

Human-equivalent hours reclaimed
(reviewer-verified)

9,794

Audited agent work items completed

120+

Reusable agent skills / playbooks

~27

External & internal tool integrations

What it is

The Platform is a **local-first, file-backed agent operating system**. Behavior is defined in version-controlled files, kept alive by always-on background services, connected to the outside world through standardized tool integrations, and steered from a single operator dashboard. Six building blocks combine into one coordinated environment:

- **Agent personas** — specialized AI "employees," each with a role, a manager, its own memory and identity. Define as many as you need.
- **Always-on services** — background processes that listen for work, run scheduled tasks, and keep the system healthy.
- **Tool integrations** — connections into the systems where work happens: messaging, calendar, tasks, content, and more.
- **Durable memory** — a knowledge base and layered memory model so agents remember decisions, people, and lessons across sessions.
- **Operator dashboard** — a single control plane to review, approve, and direct the whole team.
- **Review & automation workflows** — structured approval and self-checking loops so output is trustworthy.

OUT OF THE BOX

A pre-built core — with a signature context layer

The Platform is deliberately vanilla: **you add the agents that fit your business**. What comes pre-built is the core every agent team needs to function — most notably a **data and context management layer** that is unique to the Platform.

The Orchestrator

The always-on coordination agent. Takes in work from every channel, delegates to the right agent, enforces the chain of command, and ensures two agents never collide on the same task.

Pule — People Context

Maintains a living understanding of people, stakeholders, and relationships, so any agent acting on a task has the human context it needs.

Albert — Knowledge

Turns documents, meetings, and sources into structured, searchable knowledge the whole team can draw on.

Pule and Albert together form the Platform's differentiated data and context management capability — giving every agent long-term memory and situational awareness instead of starting each task blind. On top of this core you staff your own agents for any function — operations, marketing, research, engineering, community, support — and they inherit the same memory, coordination, governance, and observability automatically.

OPERATING LOOP

How it works

Work enters three ways: a direct human request, an external event (a message arriving in a connected channel), or a scheduled run. From there, the Orchestrator dispatches agents automatically; they coordinate through a shared control layer so nothing collides, pull in relevant context from memory and a knowledge graph, and check their own work before finishing. Anything requiring judgment is routed to a human for approval. The system is engineered to keep running through real-world disruptions — including network and corporate-proxy interruptions — by falling back across multiple AI engines rather than stalling.

What makes it unique



The defensible advantage is **integration, not any one feature**. Individually the pieces exist elsewhere; together, wired into a single reliable operating loop, they are hard to replicate:

- **Differentiated data & context management** — the Pule + Albert core makes work informed, not generic. The signature moat.
- **Durable, file-backed memory & artifacts** — knowledge and work product compound over time instead of evaporating between chats.
- **A true coordination layer** — work locks, health signals, cost tracking, and an auditable task lifecycle.
- **Resilient multi-engine dispatch** — automatic fallback across AI providers keeps agents working through outages.
- **A self-improvement loop** — the system captures its own lessons and refines its playbooks over time.
- **Governance & visibility built in** — human approval, audit trails, and a live control plane over every agent.

PROVEN RESULTS

Measured outcomes

From live operation of a reference deployment — ledger running since mid-April 2026 (under three months).

- **~3,238 human-equivalent hours reclaimed** — reviewer-verified, not self-reported.
- **9,794 audited work items** confirmed through a human review ledger with confidence-tiered scoring.
- **68+ executive briefs** produced automatically, plus recurring decision packages and weekly retrospectives.
- **Continuous, self-healing operation** — health monitoring auto-recovers failed services; humans alerted only when needed.
- **Every action logged and cost-tracked** — a complete audit and spend trail across the workforce.

SCALE & MATURITY

Production-grade

A substantial system, not a prototype.

- **120+ documented agent skills**, with a review gate for new ones.
- **~27 tool integrations** across communication, tasks, content, and internal services.
- **9+ always-on services** with automated health monitoring and self-healing.
- **A ~27-view operator dashboard** as the single human control plane.
- **A substantial custom codebase** — the dashboard alone is ~14,000 lines.
- **800+ curated knowledge notes** plus a queryable knowledge graph.

TRUST

Governance, safety & control

Trust is designed in. Sensitive actions require human approval through the dashboard or a messaging channel; a document-review cycle governs important deliverables; automated safety checks block destructive commands and secret leakage; every tool action is logged; and even adding a new agent follows a formal approval chain ending with a human owner. Agents are deliberately scoped away from out-of-bounds systems.

Vision & productization path

Today's system is **production-ready as a single-operator platform on macOS**. The natural next step — and the larger opportunity for an acquirer — is to productize it as an **enterprise-deployable, multi-tenant platform** that runs inside standard corporate environments (for example, using Microsoft Teams and Outlook for human interaction and enterprise backlog tools for task boards, with each agent holding a corporate identity).

Candid for diligence: this enterprise, cross-platform form is a clear, documented **vision — not shipped capability**. Reaching it requires additional refactoring, notably cross-platform support beyond macOS (e.g. Windows/Linux) and multi-tenant hardening.

The architectural blueprint for this path already exists, and it reuses the same proven core — the orchestration loop, the coordination layer, the Pule/Albert context management, and the tool-integration standard — on enterprise-grade, swappable components. In other words, **the hard, differentiating core is built and running; the remaining work is packaging and platform breadth**.

THE CASE

Why acquire

A working template Enterprise "AI employees" proven in daily use, with measured hours reclaimed — not a slide deck.	A defensible moat Especially the data & context management core: value lives in how the parts work together reliably — far harder to copy than any single component.	Governance-first Human-in-the-loop control and auditability that enterprises require, already embedded.	A clear runway A documented, credible path from single-operator system to multi-tenant corporate product, on the same battle-tested core.
--	---	--	--

Figures reflect a single-operator reference deployment and are reported as approximate where counts vary at runtime. Product name shown ("the Platform") is a placeholder pending final branding.